EMERGENT COMPLEXITY

Article

The Operatory Dissolution of Truth: Ontotechnics of Functional Falsehood

This essay applies, in an ontotechnical key, results established in **Ontology of the Difference Between Truth and Fiction** *(ontological base article, 28 September 2025).*

David Cota

Abstract

This essay analyses how social platforms operate as ontopolitical infrastructures whose ontotechnical design reallocates the conditions and costs of proof. By indexing value to attention capture—clicks, shares, dwell time—they privilege the performative efficacy of adherence over epistemic validity, making functional falsehood operationally superior. In this environment, filters of appearance and retentional efficacy displace confrontation with evidence, destabilizing classical accounts of truth (correspondence, coherence, pragmatism, truth as event). Empirically, statistical learning without understanding amplifies emotionally charged content and produces authorless falsity by design rather than intention. The essay proposes an operative program: friction algorithms and sustained verification time; binding multistakeholder governance (audits ex ante/ex post, objective-function transparency, model/version registries, redress); digital public goods (open provenance protocols, trust graphs, evidence repositories); and critical literacies with counter-algorithmic practices that redistribute the cost of proof. Reframing critique as intervention at the level of infrastructure, it concludes that public reason must be materially engineered so that truth regains a right it has been algorithmically denied: the right to spend time.

Keywords: ontotechnics, functional falsehood, attention economy, retentional efficacy, algorithmic amplification, dwell time, provenance protocols, trust infrastructures, chains of evidence, social epistemology, information ethics, correspondence theory, coherence theory, pragmatism, truth as event, friction algorithms, algorithmic governance, public reason, verification time, graph homophily

Social networks are not merely spaces for the circulation of utterances; more profoundly, they are ontopolitical infrastructures: techno-material devices that reconfigure the conditions of possibility for saying, hearing, and believing. Following Foucault and Deleuze, these platforms do not represent the real; they produce regimes of visibility and enunciation, operating as machines of affective governmentality. The selection of what may emerge as sayable or legitimate obeys a material logic of retention, whose criterion is not truth but the performative efficacy of adherence. This logic is inscribed in quantitative

 10.5281/zenodo.17238158

metrics—clicks, shares, dwell time—which institute a regime of value indexed to attention capture, redrawing the hierarchies of symbolic space and substituting argumentative value with retentional efficacy.

Consider the concrete case of vaccine misinformation on Facebook between 2019 and 2021. A study by Avaaz documented that the ten largest anti-vaccine pages generated 7.7 million monthly interactions, whereas the ten leading pages of public-health institutions generated 5.5 million. The false claim — "vaccines contain microchips for tracking"—circulated four times more than the WHO article "how mRNA vaccines work". The difference did not lie in epistemic quality; it lay in emotional architecture. The first formulation simultaneously activated fear, indignation, and group belonging; the second required concentration, scientific literacy, and the absence of immediate emotional gratification. The algorithm did not distinguish true from false: it identified patterns—average time of 12 seconds on the WHO post, 48 seconds on the conspiratorial post; share rate of 0.3% versus 4.7%; eight-fold higher density of emotional comments. The platform learnt a correlation: content of type A yields retention; content of type B yields abandonment. It promoted type A. Falsehood became functional not by intention, but by systemic design.

What becomes visible or viral is not what best withstands rational scrutiny, but what most effectively adheres to affects already in circulation. Language detaches from its mediating function between subject and world, becoming a matrix of emotional synchronisation. This displacement reconfigures the symbolic function: recognition shifts from the criterion of proof to that of affective resonance—that is, to the capacity for adherence. By ontotechnics I mean precisely this: the technical production of the conditions for the emergence and recognition of what may count as existent in the public sphere. This is not neutral mediation, but active configuration of what can appear and in what form. Truth, as a category of resistance to the immediate, enters this circuit at a structural disadvantage. To make the scope of this displacement explicit, it is worth clarifying the principal theories of truth at stake and how ontotechnics destabilises them.

The correspondence theory holds that a statement is true when it mirrors a state of affairs; ontotechnics does not deny the world, but it destabilises correspondence by shifting the filters of appearance and the costs of proof onto metrics of retention: what appears and persists ceases to depend on confrontation with proof and comes to depend on its circulatory compatibility.

The coherence theory understands truth as intra-systemic consistency; the algorithmic ecology tends to convert that coherence into intra-systemic closure—a ritualised self-consistency—namely, self-sustaining belief systems whose informational closure simulates validation.

The pragmatic theory reads truth as that which results from effective practices within a community of inquiry; ontotechnics perverts that outcome criterion, substituting epistemic efficacy (what resists refutation) with retentional efficacy (what maximises session time).

Truth as Event — understood here, in a materialist key, as the irruption of a new symbolic consistency that reconfigures the field — requires temporality, hesitation, and friction; the current regime rarefies the event by neutralising precisely the technical conditions that could render it possible. In this sense, ontotechnics not only subverts correspondence, coherence, and pragmatism, but also impoverishes the regime of the event, converting it into a rare exception within the digital public sphere.

Let us pick up the thread: within this informational ecosystem, validation designates, above all, fitness for circulation—not probatory confrontation. Accordingly, language is reconfigured by the

principles of circulatory efficacy. Falsehood ceases to be a moral rupture or an intentional violation of the classical epistemic contract; it becomes a structural function of the informational ecology—an effect of the algorithmic logic of visibility distribution itself.

At this point, the classical distinction between operational falsehood—that which emerges spontaneously to maintain interactive fluency, without prior calculation—and strategic falsehood—that which mobilises falsity instrumentally and in a planned manner—dissolves. Both converge in one and the same functional regime: the algorithmic optimisation of circulation. Both are assessed not by correspondence to the world, but by capacity to activate available affects and to access privileged zones of visibility.

This obliges us to reopen the question of intentionality. If falsehood is, to a large extent, a systemic effect (operational falsehood) and not merely a deliberate act (strategic falsehood), then the philosophical definition of falsehood cannot be limited to the agent's deceitful intent. Intentionality remains relevant—it distinguishes calculated deception from emergent deception—but it does not exhaust responsibility. In an ecosystem where falsity is produced by structural compatibility with retention metrics, the ethics of communication must operate on three coupled planes: (1) design responsibility—those who conceive metrics, interfaces, and promotion criteria are accountable for ontotechnical defects that generate authorless falsity; (2) institutional responsibility—news organisations, schools, platforms, and regulators that define protocols of proof, uncertainty labelling, and tempos of circulation; (3) individual responsibility—minimum duties of hesitation, elementary verification, and non-amplification where there are signs of low traceability. Intentionality thus becomes a gradient within an economy of risks: from explicit deceit to architecture-reinforced negligence. Where there is no deceit, there may be fault through uncritical adherence; where the architecture induces predictable error, there is product liability. In sum: recognising systemic falsehood does not absolve the agent; it widens the scope of imputation, shifting it from the speaker's psychologism to the engineering of the conditions of enunciation, and to distributed practices of communicative care.

The viralisation of the *Tide Pod Challenge* in January 2018 became an exemplary case. Videos of young people biting detergent capsules proliferated on YouTube and Instagram. Procter & Gamble issued notices with medical warnings about severe poisonings; poison control centres reported two hundred and twenty cases in two weeks. The institutional response was factually correct ("detergent pods are toxic and can cause oesophageal burns"), yet that true sentence generated 12,000 shares. A parodic video entitled "I ate a Tide Pod and became a superhero" garnered 2.3 million views in 48 hours. This is a paradigmatic instance of the previously defined convergence: it confirms the functional prevalence produced by systemic design.

The discursive scene is thereby transformed into a matrix of sensory feedback. What appears does so no longer by epistemic merit, but by adherence to the user's affective–cognitive profile. The public sphere ceases to be a place of discursive interaction to become a surface of continuous emotional modulation.

This continuous modulation also reconfigures the subjective experience of truth. What presents itself as "evident" comes to coincide with what has affective salience, producing a "seems-true" without the labour of proof that would sustain an "is-true". When validation is disincentivised, subjective certainty derives from the intensity of affect rather than from the resistance of the utterance, instituting a regime of instant certainty. Ethically, this shifts responsibility from the mere "do not lie" to the governance of

attentional rhythms: cultivating hesitation, tolerating delay, and suspending amplification when traceability is low. As anticipated by Jonathan Crary and Byung-Chul Han, this transformation entails the collapse of critical negativity. Falsehood ceases to be an exception—it becomes an operative rule, a systemic requirement.

The symbolic function of language undergoes a decisive reconfiguration. Understood not as mere representational coding but as a capacity for the material reorganisation of difference, that function is appropriated by a heuristic simplification characteristic of contemporary algorithmic infrastructures. This logic privileges what confirms and resonates and penalises what destabilises or complexifies. Enunciative value—that is, an utterance's material potency to gain visibility and produce effects in the public sphere—shifts from the internal complexity of the argument to its capacity for circulation.

Low-cognitive-density formulae, linear causalities, and fragments of simulated proof—decontextualised screenshots, isolated statistics—become the new effective speech acts. The performative force of the utterance is thenceforth measured by its compatibility with the algorithmic amplification system.

This symbolic reorganisation rests on a decisive material foundation: the techno-informational architecture of digital platforms. Every user action—click, abandonment, share—is recorded as operational data rather than interpretative data. Such data are not semantically understood; they are statistically correlated, as Matteo Pasquinelli stresses in his critique of the "artificial intelligence of capital". As an operator of regularities, the algorithm learns by statistical recurrence: it adjusts and reinforces without understanding. Investigations by the Mozilla Foundation and Harvard University documented that users who watched videos on vegetarian diets were progressively directed towards extreme veganism, then radical anti-speciesism, and finally conspiracy theories about the food industry. The algorithm did not understand ideological positions; it identified a pattern: users who watched A spent longer on B, and longer still on C. YouTube explained that the algorithm optimises session time and had empirically observed that increasingly radical content retains attention more effectively. The platform promoted radicalisation not out of ideological conviction but through economic optimisation: heightened emotional intensity generated dwell time, and dwell time generated advertising revenue. The statistical correlation—radicality → retention—sufficed; the semantic understanding of what "radicality" is, and of its social effects, was irrelevant.

Without lapsing into essentialisms of technology, the analysis must also open onto the philosophy of technology. In Bernard Stiegler, technique appears as *pharmakon*—simultaneously poison and remedy—because it externalises memory and retains the symbolic, reconfiguring our circuits of attention and of credit in the common; in this key, digital ontotechnics institutes a pharmacology of retention, where any cure can only come through the redesign of attentional prostheses and of the regime of retentions. Don Ihde shows that every technique mediates perception and action across multiple scales (it amplifies, reduces, translates), which allows us to read algorithms not as passive filters but as structures of co-intentionality that redistribute agency among users, metrics, and interfaces. Albert Borgmann distinguishes devices that conceal effort and maximise convenience from the focal paradigm, which summons attentive practice; contemporary platforms operate as devices of cognitive convenience that dissolve probatory friction—hence a politics of truth should refocus the public sphere by means of rhythms, formats, and rituals that reinstate attention as a shared practice.

It bears emphasising that this algorithmic configuration is neither neutral nor inevitable. The choice to optimise session time at the expense of truthfulness is a conscious business decision, determined by the advertising-revenue structure of the platforms. Internal documents made public in recent years—most notably the *Facebook Files* disclosed by Frances Haugen in 2021—show that these companies possess detailed data on the polarising, addictive, and disinformative effects of their recommendation systems. Meta knew, at least since 2018, that its algorithm privileged content generating "significant anger" because it maximised engagement. YouTube had likewise identified, in internal investigations in 2019, that automatic playback progressively directed users towards more extreme content. The maintenance of this design does not stem from technical ignorance or operational impossibility, but from economic calculation: functional falsehood is more profitable than truth. Each additional second of retention translates into measurable advertising exposure, and each engagement metric underwrites these corporations' market valuation. Technically, we have already noted the statistical blindness; politically, the agency is human. The ontotechnics of functional falsehood is thus inseparable from a political economy of attention in which cognitive capture constitutes the business model. What presents itself as systemic necessity is, in fact, a strategic choice—reversible, modifiable, yet deliberately maintained so long as value extraction depends upon it.

In this new regime, the public discursive space ceases to function as an arena of shared justification, and comes to operate as a laboratory of affective adherence. Validation becomes a structurally disincentivised act—not through explicit censorship, but because the infrastructure itself shifts the cognitive effort of verification towards immediate gratification. The cost of proof is externalised: it is no longer jointly guaranteed between senders and receivers within a communicative contract, but absorbed and dissolved by the technical system of visibility. Proof thus re-encounters the previously noted cost differential (*time for proof* versus fluency), and is penalised for interrupting the flow of attention.

The language that seeks to maintain a bond with truth thus finds itself compelled to incorporate the codes of viral fiction—narratives of rapid digestion, affective condensation, and immediate visual appeal. This adaptation, however, undermines the critical potency of truthful language. When truth is forced to mimetise the expressive modes of falsehood in order to compete for attention, it compromises hesitation, abdicates openness to refutation, and suppresses the time required for verification. What might have functioned as a survival strategy becomes, through cumulative effect, an operative dissolution of the discourse's critical function.

It is worth invoking, by way of counterpoint, the tradition that conceives language as a space of passage and resistance. From the Socratic gesture of public interrogation to Derrida's proposal of deconstruction, language has been thought as a place where truth is not given but sought—constructed in the tension between the uttered and the un-said, between what is instituted and what exceeds it. Such a conception requires temporality, hesitation, and openness to dissonance: precisely what the current technical environment systematically neutralises. The speed of the feed is not merely a rhythm—it is an organising principle that prevents the emergence of truth as event. In the light of phenomenology, truth may be conceived as a modality of appearing: not merely propositional correctness, but a situated unconcealment (*aletheia*) that requires temporal thickness and the body's lingering attention. In Merleau-Ponty, embodied perception founds horizons of meaning that are not neutral: one's own body functions as a matrix for selecting what may figure as evident. In an ontotechnical environment of low cognitive density,

this grammar of appearing is compressed: the phenomenological window required for the true to show itself as such is narrowed. Hermeneutics helps to name what is lost: the hermeneutic circle (Gadamer)—between pre-understanding and confrontation with text/world—demands rhythms of back-and-forth that retention metrics shorten. In Ricœur, the passage through narrative and action institutes interpretative distances that allow the revision of judgement; the feed suppresses that distance, binding evidence to immediacy. Phenomenology and hermeneutics thus converge on an operative point: without a rite of attention and distancing, the subjective experience of the true is depotentiated by a design that favours rapid affective synthesis.

At this point a material redesign of the conditions of enunciation becomes unavoidable. This diagnosis requires a theoretical displacement: from the moral critique of falsehood to the ontotechnical analysis of the conditions of enunciation. The issue is not to denounce false contents, but to understand the material devices that render them functionally superior. It is not enough to try to make truth competitive with the very instruments that erase it. From here follows the explicit redistribution of the cost of verification, via infrastructural measures. Recognising that platforms do not self-regulate—given that functional falsehood underwrites their business model—jurisdictions such as the European Union and Australia are beginning to impose regulatory frameworks that require algorithmic transparency, accountability for amplified content, and mechanisms for the mitigation of misinformation. The Digital Services Act (DSA), adopted in 2022, represents a first step, by obliging very large platforms to undergo external audits of their recommender systems and risk assessments regarding systemic effects. Although insufficient—since they remain chiefly oriented towards the moderation of individual items of content rather than the transformation of the logic of amplification—these instruments demonstrate that regulatory intervention in technical architecture is politically viable. From this diagnosis there follows an operative criterion for institutions.

In Floridi, the *infosphere* is not a mere repository but a common ontological environment: to act technically is to intervene in the informational fabric. From this follow positive duties—to preserve, enrich, and not degrade informational value—which, in our context, translate into obligations of design (metrics, interfaces, promotion policies) and governance (audits, traceability, uncertainty labelling). On the side of social epistemology, knowledge is a good co-produced by practices of testimony, expertise, and public trust; validity is not only an attribute of the utterance but a relational property of networks and institutions. This demands trust infrastructures (verifiable procedures, source records, legible chains of evidence) and a response to credibility asymmetries that amplify the false (echo chambers, authority biases) and silence the true (epistemic injustice). In sum: information ethics provides the normative standard (do not degrade the *infosphere*; raise its value) and social epistemology specifies the collective mechanisms of validation (organisation of testimony, distribution of trust, responsibility by design); both converge in our criterion: without networks of proof and trust built into the architecture, truth loses the world. Such an exigency entails regulation that does not confine itself to penalising false content *after the fact*, but that mandates the redesign of algorithmic success metrics, privileges time for reflection over speed of reaction, and shifts the burden of proof from consumers to publishers and amplifiers.

We can therefore spell out, by way of example, principles for a normativity of infrastructure: (i) goal-oriented transparency—not only concerning moderation processes, but concerning the objective function of recommender systems themselves (with declared weights for truthfulness, diversity, and well-

being, and aggregated public reports); (ii) independent, continuous auditability, enabled by audit APIs and sandboxes with synthetic data for third-party testing, with duties to correct when foreseeable disinformative effects are detected; (iii) normative prioritisation of truthfulness over engagement, instantiated in multi-criteria optimisation, with explicit penalties for patterns associated with low traceability and breaks in the chain of evidence; (iv) promotion of diversity of perspectives, via diversity injectors and limits on graph homophily (ensuring minimum exposure to verified independent sources within defined time windows); (v) responsibility for social impact, with public risk assessments, provenance records, and chains of custody for amplified content.

Operationally, this translates into friction algorithms: deliberate latencies on viral shares, "read before sharing" prompts, dynamic re-amplification limits, the deceleration of low-verifiability trends, pause quotas that protect time for proof, and uncertainty labels that invite public hesitation. All of this should converge on a new success metric: sustained verification time, not merely passive dwell.

Models of algorithmic governance. Beyond the Digital Services Act (DSA), a binding multi-stakeholder model is required, one that involves public regulators, civil society, academia, and companies in the design, testing, and review of platforms. In addition to the audit APIs already mentioned, this entails: (i) mandatory ethical audits *ex ante* (before launching/altering recommender systems) and *ex post* (on real-world data), with publication of summary reports; (ii) independent authorities with a technical mandate to impose design corrections when foreseeable effects of misinformation, polarisation, or discrimination are detected (injunctive and fining powers); (iii) public records of model versions, objective-function parameters, and catalogues of training and evaluation data (with privacy safeguards), enabling decision traceability; (iv) redress mechanisms for parties harmed by algorithmic decisions, with deadlines and duties to respond; (v) regulatory sandboxes for supervised experimentation with truthfulness and diversity metrics.

Role of the individual and education for resistance. Resistance is not exhausted by regulation or institutional design: it requires critical media and digital literacy that combines provenance reading skills (tracing sources, recognising legible chains of evidence and gaps), functional understanding of algorithms (knowing, in operational terms, what objective function, training signal, and retention metrics are, and their selection effects), and an ethics of attention (training in delay, suspension of amplification, tolerance for dissonance). Pedagogically, this implies curricula that teach how to reconstitute content trajectories (from capture to *ranking*), simulate recommendations to show how small interactional variations produce informational drift, and practise verification protocols with explicit proof-times. Civically, it implies public rites of proof—distributed verification labs, slow-reading circles, community devices for uncertainty labelling—that restore to truth its shared time. Thus, "public reason" ceases to be a normative abstraction and becomes a trainable competence: organising deep attention under conditions of noise, sustaining hesitation where the infrastructure accelerates, and maintaining operative coherence under the pressure of mere circulatory fitness.

Resistance and Counter-Algorithmics. "Socratic resistance" becomes effective in the digital only when public interrogation is translated into material procedures: auditable lists of premises (with records of algorithmic decision-making), a right of interpellation embedded in interfaces (clear access to the objective function, to data provenance, and to operative explanations), and mechanisms for adversarial rebuttal (prioritised counter-statements with a chain of evidence). "Derridean deconstruction" here shifts

to design: exposing and undoing binaries and hierarchies embedded in metrics (for example, engagement > truthfulness; salience > proof), introducing temporal difference (pauses, delays, re-amplification limits) and structural difference (forced diversification of sources) as techniques of denaturalisation. To resist ceases to be merely content critique and becomes a practice of engineering and technological activism: browser extensions that restore provenance to view, extensions flagging low traceability, alternative "slow" feeds with explicit weights for truthfulness and diversity, and civic hackathons for bias tests and public stress tests of platforms. This counter-algorithmics does not aim to paralyse circulation, but to redistribute the cost of proof and reopen the time in which truth can resist—prolonging, in a materialist key, Socratic interrogation and Derridean deconstruction at the very level of infrastructures.

As an alternative to the current business model, we should institute digital public goods or common digital infrastructures: open protocols of provenance and content chains of custody; public indices of informational quality and auditable trust graphs; interoperable evidence repositories for journalism and science; public APIs providing access to civic metrics (verification time, source diversity) and interoperability obligations between platforms. These arrangements reorient the ecosystem towards the public interest, shifting the dominant rationality from attention extraction to the shared production of conditions of proof.

Examples of counter-algorithmic institutions. A number of existing models already test, with differing degrees of maturity, the principles defended here.

Federated networks (Mastodon/ActivityPub). They avoid a single centre of decision and reduce reliance on engagement metrics (chronological/localised timelines by default; instance-level public moderation rules). Success: greater transparency and the possibility of choosing norms. Challenges: fragmentation, cross-instance content discovery, distributed moderation, and economic sustainability.

Knowledge platforms (Wikipedia). They practise legible validation (verifiability policies, edit histories, talk pages, auditable reversions). Success: public proof and reversibility. Challenges: systemic bias, targeted harassment, and reliance on volunteers.

"Slow" journalism and civic investigations (*Tortoise, Zetland, De Correspondent; ProPublica, The Markup*). They prioritise time for proof over speed, with membership models and the publication of evidence bases. Success: depth and traceability. Challenges: scale and appeal in an ecosystem habituated to rapid gratification.

Fact-checking ecosystems (networks such as IFCN and national initiatives). They introduce chains of evidence and uncertainty labels. Success: public corrections and minimum standards. Challenges: latency and limited reach without platform-level integration.

Open annotation layers (*Hypothes.is*). They add visible adversarial commentary to content. Challenges: adoption inertia.

Taken together, these experiments show that slow infrastructures and legible validation formats are possible: they gain in proof and public accountability what they lose in speed and scale. The next step is to institutionalise these mechanisms—interoperable, auditable, and economically sustainable—so that they no longer depend on heroic exceptions.

What is at stake is not merely the circulation of content, but the very possibility of a public reason as a material space for the reorganisation of meaning. Functional falsehood is neither a moral accident nor an epistemic pathology: it is a necessary effect of an algorithmic ecology designed to reward what

adheres, not what is verified. Critique cannot be confined to normative diagnoses: it must become an intervention on the ontotechnical plane already specified, requiring an ethics of infrastructure—an ethics that does not merely judge the utterance, but redesigns the material conditions of its emergence. Only thus will it be possible to restore to language its most radical power: to resist that which merely works.

Truth requires counter-algorithmic institutions oriented by this design criterion.

Justice in the public sphere begins by returning to truth the right to spend time.

References

- Borgmann, Albert. 1984. *Technology and the Character of Contemporary Life: A Philosophical Inquiry*. Chicago: University of Chicago Press.
- Crary, Jonathan. 2013. *24/7: Late Capitalism and the Ends of Sleep*. London: Verso.
- Deleuze, Gilles. 1992. "Postscript on the Societies of Control." *October* 59: 3–7.
- European Union. 2022. *Regulation (EU) 2022/2065 of the European Parliament and of the Council of 19 October 2022 on a Single Market for Digital Services (Digital Services Act)*. *Official Journal of the European Union* L 277 (27 October 2022). https://eur-lex.europa.eu/eli/reg/2022/2065/oj
- Floridi, Luciano. 2011. *The Philosophy of Information*. Oxford: Oxford University Press.
- ———. 2013. *The Ethics of Information*. Oxford: Oxford University Press.
- Foucault, Michel. 2007. *Security, Territory, Population: Lectures at the Collège de France, 1977–78*. Edited by Michel Senellart. Translated by Graham Burchell. Basingstoke: Palgrave Macmillan.
- Gadamer, Hans-Georg. 2004. *Truth and Method*. 2nd, revised ed. Translated by Joel Weinsheimer and Donald G. Marshall. London and New York: Continuum.
- Han, Byung-Chul. 2017. *In the Swarm: Digital Prospects*. Cambridge, MA: MIT Press.
- Ihde, Don. 1990. *Technology and the Lifeworld: From Garden to Earth*. Bloomington: Indiana University Press.
- Merleau-Ponty, Maurice. 2012. *Phenomenology of Perception*. Translated by Donald A. Landes. London and New York: Routledge.
- Mozilla Foundation. 2021. "*YouTube Regrets*: What We Learned by Tracking YouTube's Recommendation System." Mozilla Foundation (Report). https://foundation.mozilla.org
- Ribeiro, Manoel H., Raphael Ottoni, Robert West, Virgílio A. F. Almeida, and Wagner Meira Jr. 2020. "Auditing Radicalization Pathways on YouTube." *arXiv* preprint arXiv:1908.08313. https://arxiv.org/abs/1908.08313
- Ricoeur, Paul. 1984. *Time and Narrative: Volume 1*. Chicago: University of Chicago Press.

- Stiegler, Bernard. 2010. *Taking Care of Youth and the Generations*. Stanford, CA: Stanford University Press.
- ———. 2013. *What Makes Life Worth Living: On Pharmacology*. Cambridge: Polity.
- The Wall Street Journal. 2021. "The Facebook Files." September–October 2021 (series page). https://www.wsj.com/topics/series/the-facebook-files
- U.S. Senate Committee on Commerce, Science, and Transportation, Subcommittee on Consumer Protection, Product Safety, and Data Security. 2021. "Testimony of Frances Haugen." 5 October 2021. https://www.commerce.senate.gov
- W3C. 2018. "ActivityPub." W3C Recommendation, 23 January 2018. https://www.w3.org/TR/activitypub/
- International Fact-Checking Network (IFCN). 2022. "IFCN Code of Principles." Poynter Institute. https://ifcncodeofprinciples.poynter.org
- Hypothesis. 2023. "Collaborate & Annotate with Hypothesis." https://web.hypothes.is
- Avaaz. 2020. *Facebook's Algorithm: A Major Threat to Public Health*. 19 August 2020. https://secure.avaaz.org/campaign/en/facebook_threat_health/